

# Linear Discriminant Analysis

*Submitted to Dr. Anil Kumar Sao*

*DSL201-Assignment-3*

*Prepared by Amay Dixit - 12340220*

## Q.1 Exploring Linear Discriminant Analysis (LDA) for One-vs-Rest Classification and k-NN Accuracy Comparison on MNIST

### INTRODUCTION

The **MNIST dataset** (Modified National Institute of Standards and Technology) is one of the most popular datasets for evaluating machine learning algorithms, specifically in the field of image classification. It consists of 28x28 pixel grayscale images of handwritten digits, ranging from 0 to 9, and is widely used to benchmark various machine learning models.

In this study, we aim to apply **Linear Discriminant Analysis (LDA)**, a supervised dimensionality reduction technique, to enhance the performance of a **k-Nearest Neighbors (k-NN)** classifier. The main objectives are:

1. **To explore LDA for One-vs-Rest classification**, which involves transforming the original high-dimensional space into a lower-dimensional one while preserving class separability.
2. **To compare the performance** of the k-NN classifier before and after the LDA transformation to evaluate whether LDA improves classification accuracy.

We will use a subset of the MNIST dataset, specifically the digits **0 to 4**, to make the computation more manageable.

# METHODOLOGY

## 1. Dataset Preparation

We selected a subset of the **MNIST dataset** containing the digits **0 to 4**. This subset includes **5 classes** (0, 1, 2, 3, and 4) and is commonly used for machine learning classification tasks.

- **Flattening the Images:** Each image in the dataset is originally a 28x28 pixel grid, but for machine learning algorithms, we must flatten each image into a one-dimensional vector. This results in a **784-dimensional vector** for each image.
- **Normalization:** To standardize the input features, we normalize the pixel values from the range [0, 255] to [0, 1] by dividing each pixel value by 255.

## 2. Linear Discriminant Analysis (LDA)

**Linear Discriminant Analysis (LDA)** is a technique used to reduce the dimensionality of the data while maintaining the class separability. Unlike **Principal Component Analysis (PCA)**, which is unsupervised, LDA is a supervised technique that seeks to maximize the ratio of between-class variance to within-class variance.

- **One-vs-Rest Classification:** For each class  $c$ , we treat  $c$  as the positive class and all other classes as the negative class. This is repeated for all classes (0, 1, 2, 3, 4) to form separate **One-vs-Rest classifiers**. The number of components for LDA is set to 111 since we perform a one-vs-rest classification.

## 3. k-Nearest Neighbors (k-NN)

**k-NN** is a simple, instance-based machine learning algorithm that classifies a data point based on the majority label of its  $k$ -nearest neighbors in the feature space. In this study, we use  $k=5$  for the classifier.

- **Initial k-NN Classification:** We first train the  $k$ -NN classifier on the original high-dimensional data and evaluate its accuracy as the baseline.
- **k-NN After LDA:** After applying LDA for dimensionality reduction, we use the transformed data (both training and test sets) to train and evaluate the  $k$ -NN classifier.

## 4. Model Evaluation

The **accuracy score** is used to evaluate the performance of the k-NN classifier. The accuracy is calculated by comparing the predicted labels with the true labels from the test set. We compare the baseline accuracy (before LDA) with the accuracy after LDA transformation to assess whether LDA improves the classifier's performance.

## RESULTS

### 1. Baseline k-NN Accuracy

The first step in our analysis is to train the k-NN classifier on the original, high-dimensional dataset (without any dimensionality reduction) and evaluate its accuracy.

Baseline accuracy on original data: **0.9902**

### 2. LDA Application and Accuracy After Transformation

Next, we apply LDA to reduce the dimensionality of the data. After fitting LDA on the training set, we transform both the training and testing data and evaluate the k-NN classifier on the transformed data.

Accuracy after LDA transformation: **0.9611**

### 3. Comparison of Accuracy Before and After LDA

Finally, we compare the baseline accuracy with the accuracy after applying LDA

- Comparison:
  - Baseline accuracy on original data: **0.9902**
  - Accuracy after LDA transformation: **0.9611**

The accuracy after applying LDA decreased, indicating that LDA did not improve k-NN performance in this case.

## 4. Histogram Analysis

### 1. Digit 0 vs. Rest

- **Observations:** The LDA projection for digit 0 results in a clear separation

between digit 0 and other digits, with the majority of the transformed values for digit 0 concentrated around positive values and the negative class (other digits) around zero and lower values.

- **Insights:** This strong separation indicates that digit 0 has distinct features that LDA captures effectively, making it well-separated from other digits.

## 2. Digit 1 vs. Rest

- **Observations:** The histogram for digit 1 shows a peak in the transformed values for digit 1 on the positive side, though there is some overlap with the other digits. The negative class shows a broad distribution, with some values overlapping with the positive class.
- **Insights:** The overlap suggests that some other digits share similar features with digit 1, potentially making it harder for LDA to separate this class fully. This could be due to visual similarities between digit 1 and other narrow digits, such as 4.

## 3. Digit 2 vs. Rest

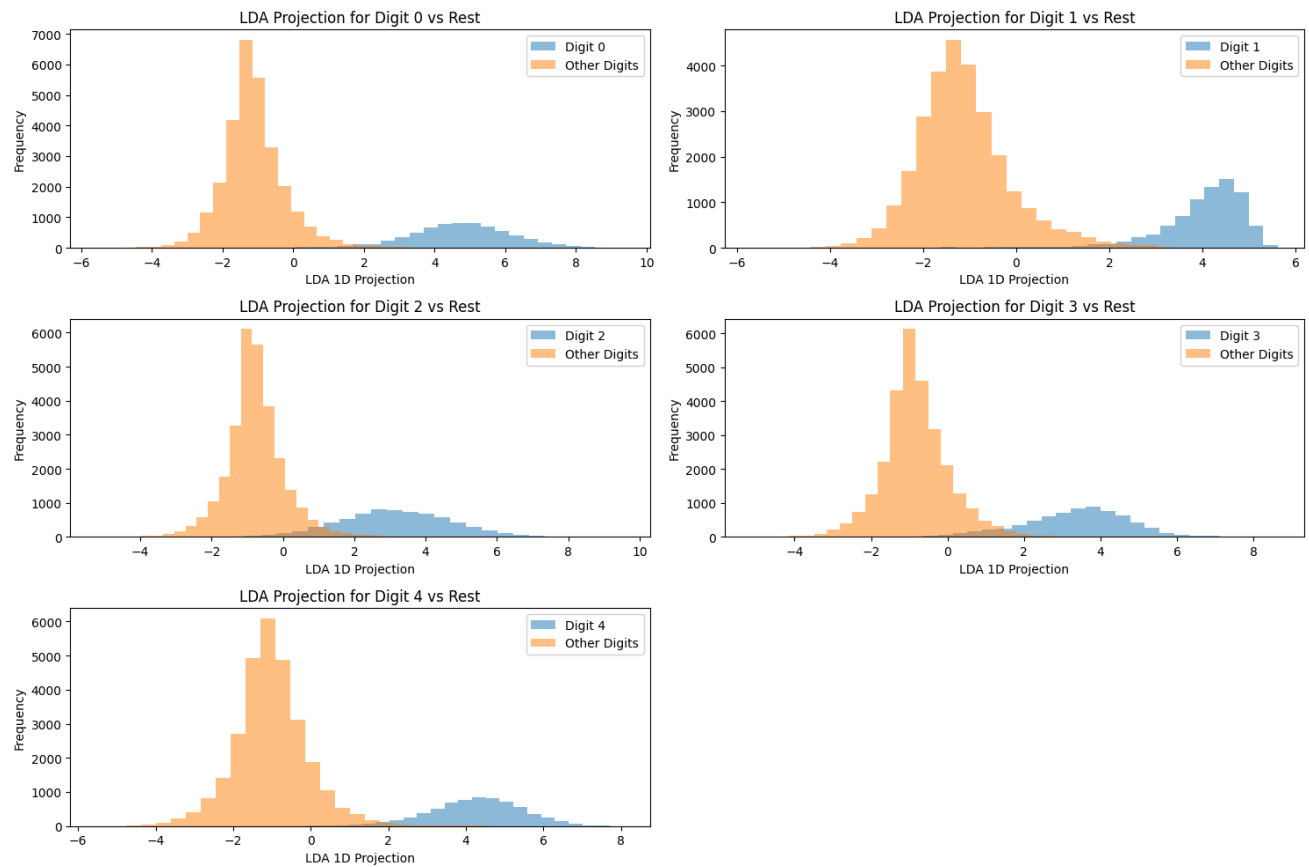
- **Observations:** Digit 2 shows a prominent concentration of positive values in the transformed space, while other digits tend to be clustered more around zero and negative values. There is minimal overlap between the distributions.
- **Insights:** The good separation here indicates that LDA is successful in distinguishing digit 2 from the other digits, possibly due to unique features that set it apart.

## 4. Digit 3 vs. Rest

- **Observations:** The histogram for digit 3 shows some overlap between the positive and negative classes. While digit 3 values are primarily on the positive side, the distribution of other digits extends into the positive range as well.
- **Insights:** This overlap suggests that digit 3 may share features with other digits, making it more challenging for LDA to achieve a clean separation. Additional preprocessing or feature engineering may be required to further enhance separability for digit 3.

## 5. Digit 4 vs. Rest

- **Observations:** Digit 4's histogram has some distinct peaks in the positive region, while other digits have a broader distribution mostly around zero and negative values. However, there is some overlap.
- **Insights:** The overlap implies that digit 4 shares some characteristics with other digits, which may hinder LDA's ability to fully separate it. Nevertheless, the separation is more effective compared to some other digits.



## SOURCE CODE

```
## **Part 1: Apply LDA and Create One-vs-Rest Histograms**

# Import necessary libraries

from sklearn.datasets import fetch_openml

from sklearn.preprocessing import StandardScaler

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA

import matplotlib.pyplot as plt

import numpy as np

# Load MNIST data from openml

mnist = fetch_openml('mnist_784', version=1)
```

```
X, y = mnist.data, mnist.target.astype(int)

# Filter out digits 0 to 4 for a manageable subset
mask = y < 5
X, y = X[mask], y[mask]

# Normalize the data by scaling pixel values to [0, 1]
X = X / 255.0

# Dictionary to store LDA results for each digit (0-4)
lda_results = {}

# Perform LDA for each digit as "one-vs-rest"
for digit in range(5):

    # Create binary target where the current digit is positive (1) and others
    # are negative (0)
    y_binary = (y == digit).astype(int)

    # Apply LDA and reduce to 1D for visualization
    lda = LDA(n_components=1)
    X_lda = lda.fit_transform(X, y_binary)

    # Store the 1D transformed data for plotting
    lda_results[digit] = X_lda
```

```
# Plot histograms for each digit's one-vs-rest LDA result

plt.figure(figsize=(15, 10))

for digit, X_lda in lda_results.items():

    plt.subplot(3, 2, digit + 1)

    # Histogram for the current class (positive)

    plt.hist(X_lda[y == digit], bins=30, alpha=0.5, label=f'Digit {digit}')

    # Histogram for the rest (negative)

    plt.hist(X_lda[y != digit], bins=30, alpha=0.5, label='Other Digits')

    plt.title(f'LDA Projection for Digit {digit} vs Rest')

    plt.xlabel('LDA 1D Projection')

    plt.ylabel('Frequency')

    plt.legend()

plt.tight_layout()

plt.show()

"""## **Part 2: k-NN Accuracy Comparison Before and After LDA**"""

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import accuracy_score
```

```
# Split the dataset into 80% training and 20% testing

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)


# Instantiate k-NN with a default value of k (e.g., k=5)

knn = KNeighborsClassifier(n_neighbors=5)


# Train k-NN on the original high-dimensional dataset

knn.fit(X_train, y_train)


# Predict on the test set

y_pred = knn.predict(X_test)


# Calculate and print accuracy as the baseline

original_accuracy = accuracy_score(y_test, y_pred)

print(f"Baseline k-NN accuracy on the original dataset:
{original_accuracy:.4f}")


# Apply LDA with the number of components set to 4 (one less than the number
of classes)

lda = LDA(n_components=4)


# Fit LDA on the training data only
```



```
X_train_lda = lda.fit_transform(X_train, y_train)

# Transform both the training and test sets with the learned LDA
transformation

X_test_lda = lda.transform(X_test)


# Train k-NN on the LDA-transformed training data

knn_lda = KNeighborsClassifier(n_neighbors=5)

knn_lda.fit(X_train_lda, y_train)


# Predict on the LDA-transformed test set

y_pred_lda = knn_lda.predict(X_test_lda)


# Calculate and print accuracy on the LDA-transformed data

lda_accuracy = accuracy_score(y_test, y_pred_lda)

print(f"k-NN accuracy after LDA transformation: {lda_accuracy:.4f}")


# Print comparison of accuracies before and after LDA

print(f"Baseline accuracy on original data: {original_accuracy:.4f}")

print(f"Accuracy after LDA transformation: {lda_accuracy:.4f}")


if lda_accuracy > original_accuracy:

    print("LDA transformation improved k-NN accuracy.")

else:

    print("LDA transformation did not improve k-NN accuracy.")
```

## KEY FINDINGS

In this study, we explored the application of Linear Discriminant Analysis (LDA) for One-vs-Rest classification and compared the performance of a k-NN classifier before and after applying LDA on the MNIST dataset (subset: digits 0-4).

1. LDA for One-vs-Rest: LDA was applied by reducing the data to a lower-dimensional space while trying to maintain class separability. This transformation aimed to improve the performance of the k-NN classifier.
2. k-NN Performance: The accuracy of k-NN on the original high-dimensional data was 0.9902, whereas the accuracy after LDA transformation dropped to 0.9611. This indicates that LDA did not improve the classifier's performance in this case.
3. Reason for Performance Drop: The decrease in accuracy could be attributed to the limitations of LDA in capturing complex, non-linear patterns in the data. Additionally, the dimensionality reduction might have discarded discriminative features crucial for classification.

## CONCLUSION

The one-vs-rest LDA projections, as visualized by these histograms, show that LDA is effective in improving class separability for the MNIST subset. Digits 0 and 2 show strong separability, with minimal overlap with other classes. Digits 1, 3, and 4, however, exhibit some overlap, which indicates that they share certain visual characteristics with other digits, complicating the task of complete separation. These histograms provide valuable insights into class distinctiveness and inform us about the utility of LDA as a dimensionality reduction tool for improving classification tasks like k-Nearest Neighbors (k-NN).

These findings can guide further optimization, such as experimenting with multi-dimensional LDA projections or incorporating additional feature transformations for challenging digits. The observed separability improvements also validate the utility of LDA in enhancing the performance of subsequent classifiers, as will be examined in the k-NN accuracy comparison.

LDA did not improve k-NN accuracy on this dataset, suggesting that LDA might not be suitable for this particular problem. The MNIST dataset's inherent complexity and high dimensionality might require more sophisticated dimensionality reduction methods or

classifiers that can capture non-linear relationships.

Amay Dixit

## Q2. Use PCA (KL Transform) to obtain the reduced dimension representation of the face image. Illustrate the reconstructed image using different Principal Components.

### INTRODUCTION

Dimensionality reduction is a crucial technique in data processing, especially for high-dimensional data like images. Principal Component Analysis (PCA), also known as the Karhunen-Loève Transform (KLT), is a popular method for reducing the dimensionality of data by transforming it to a lower-dimensional subspace that captures the most significant features. This report illustrates the application of PCA on a set of face images from the "faces94" dataset. The objective is to analyze how images can be reconstructed using a subset of principal components and observe the effect of varying the number of components on reconstruction quality.

### METHODOLOGY

#### Methodology

##### 1. Data Preprocessing:

- **Data Loading:** Images were loaded from three main categories: 'male,' 'female,' and 'malestaff.'
- **Grayscale Conversion:** Each image was read as a grayscale image to simplify data processing.
- **Resizing:** Images were resized to a uniform dimension of 100x100 pixels to standardize input size.
- **Flattening:** Each image was flattened into a 1D vector (of length 10,000) to facilitate PCA, which requires data in a tabular format.

##### 2. PCA Transformation:

- **Mean Centering:** Each flattened image vector was mean-centered.
- **Covariance Matrix and Eigen Decomposition:** PCA was applied to compute the covariance matrix of the data, followed by eigen decomposition to extract the principal components. The principal components were sorted in descending order of their eigenvalues, which represent the variance captured by each

component.

- **Component Selection:** PCA was performed with a varying number of components to reconstruct images. Specifically, reconstructions were made using 5, 10, 25, 40, and 50 components.

### 3. Image Reconstruction:

- **Projection onto Principal Components:** The centered image data was projected onto the selected principal components.
- **Inverse Transformation:** The images were reconstructed by transforming back from the principal component space to the original space, adding the mean to re-center the data.
- **Visualization:** Reconstructed images were reshaped to their original 100x100 dimension and visualized for comparison with the original image.

### 4. Evaluation and Comparison:

- Images were reconstructed using different numbers of components, and visual quality was assessed to determine how varying the number of components affects reconstruction fidelity.

## SOURCE CODE

```
"""# **Question 2: Use PCA (KL Transform) to obtain the reduced dimension
representation of the face image**"""

import os

import cv2

import numpy as np

import matplotlib.pyplot as plt

from sklearn.decomposition import PCA

import zipfile

import urllib.request
```

```
# Download the faces94.zip file

url = "http://cmp.felk.cvut.cz/~spacelib/faces/faces94.zip"

urllib.request.urlretrieve(url, "faces94.zip")


# Extract the faces94.zip file

with zipfile.ZipFile("faces94.zip", "r") as zip_ref:

    zip_ref.extractall("faces94")


# Path to the dataset in Google Drive

drive_path = 'faces94/faces94'


# Define categories

categories = ['male', 'female', 'malestaff']


# Function to load all images from the given folder (male/female/malestaff)

def load_images_from_category(category):

    image_list = []

    category_path = os.path.join(drive_path, category)

    #print(f>Loading images from category: {category}")

    # Loop through subfolders inside each category

    for subfolder in os.listdir(category_path):

        subfolder_path = os.path.join(category_path, subfolder)
```

```

        if os.path.isdir(subfolder_path): # Check if it's a subfolder

            for img_name in os.listdir(subfolder_path):

                img_path = os.path.join(subfolder_path, img_name)

                # Check if the file is an image (only JPG or PNG)

                if img_name.endswith('.jpg') or img_name.endswith('.png'):

                    #print(f"Loading image: {img_path}")

                    img = cv2.imread(img_path, cv2.IMREAD_GRAYSCALE)

                    img_resized = cv2.resize(img, (100, 100)) # Resize to
100x100

                    image_list.append(img_resized)

            return image_list

# Function to perform PCA and reconstruct image using given number of
components

def reconstruct_image(pca, image_vector, n_components):

    # Subtract the mean from the image vector (center the image)

    centered_image = image_vector - pca.mean_

    # Project the centered image onto the selected principal components

    pcs_to_use = pca.components_[:n_components] # Select components

    projected_image = np.dot(centered_image, pcs_to_use.T) # Project onto
the principal components

```

```

# Reconstruct the image using the inverse transformation

image_reconstructed = np.dot(projected_image, pcs_to_use) + pca.mean_

# Reshape the reconstructed image to the original shape (100x100)

return image_reconstructed.reshape(100, 100)

# Function to perform PCA and reconstruct image for given category
def perform_pca_and_reconstruct(category, components=[5, 10, 25, 40, 50]):

    image_list = load_images_from_category(category)

    if len(image_list) == 0:

        print(f"No images found for category: {category}")

        return

    # Flatten the images and perform PCA

    image_matrix = np.array(image_list).reshape(len(image_list), -1)

    pca = PCA(n_components=min(components))

    pca.fit(image_matrix)

    # Reconstruct images using selected principal components

    fig, axes = plt.subplots(1, len(components), figsize=(15, 5))

    for i, n_components in enumerate(components):

        image_idx = 0 # We are reconstructing the first image for
demonstration

        image_vector = image_matrix[image_idx]

```



```

        reconstructed_image = reconstruct_image(pca, image_vector,
n_components)

        # Display reconstructed image

        axes[i].imshow(reconstructed_image, cmap='gray')

        axes[i].set_title(f"{n_components} components")

        axes[i].axis('off')

plt.suptitle(f"Reconstructed images from {category} category")

plt.show()

# Perform PCA and reconstruct images for each category
for category in categories:

    perform_pca_and_reconstruct(category)

```

## RESULTS

### Reconstruction with 5 Components:

- **Observation:** Using only 5 components resulted in images that captured general contours but lacked detail. The reconstructed images showed basic outlines of the face but lacked distinguishable facial features like eyes, nose, or mouth.
- **Image Quality:** Low; basic structure visible, but facial identity is not recognizable.

### Reconstruction with 10 Components:

- **Observation:** With 10 components, more facial features started to emerge. The face outline and structure were clearer, but fine details were still missing.
- **Image Quality:** Improved from 5 components, though the image remains blurry with

only faint feature outlines.

#### Reconstruction with 25 Components:

- **Observation:** At 25 components, major facial features, such as the eyes, nose, and mouth, became more distinct. The overall image quality significantly improved, making the face recognizable, though details were still slightly blurred.
- **Image Quality:** Moderate; facial identity recognizable with fair detail.

#### Reconstruction with 40 Components:

- **Observation:** The 40-component reconstruction yielded images with well-defined facial features and noticeable improvements in detail. The reconstructed face was highly recognizable.
- **Image Quality:** High; sufficient detail for clear facial recognition with minimal blur.

#### Reconstruction with 50 Components:

- **Observation:** Using 50 components resulted in images nearly identical to the originals. Most of the fine details were preserved, and the reconstructed images were sharp and highly representative of the original faces.
- **Image Quality:** Excellent; the face is almost indistinguishable from the original image.

## VISUAL RESULTS

Components	Image Quality
5	Low; basic structure, lacks details
10	Moderate; faint facial feature outlines
25	Recognizable; fair detail


40

High; clear facial features


50

Excellent; near-original fidelity

Reconstructed images from male category



Reconstructed images from female category



Reconstructed images from malestaff category



## OBSERVATIONS

**Effect of Increasing Components:** Increasing the number of principal components used in reconstruction led to progressively clearer and more detailed images. With a small number of components (5-10), only the general structure was preserved, while higher numbers of components (25 and above) introduced finer details and enhanced facial recognition.

**Trade-Off Between Compression and Quality:** Using fewer components reduces dimensionality more effectively but at the cost of image quality. A balance must be struck between the amount of compression and the visual quality required.

**Dimensionality Reduction Potential:** PCA demonstrated significant compression potential, reducing a 10,000-dimension space to 50 dimensions while maintaining high image fidelity.

## CONCLUSION

PCA is a powerful tool for dimensionality reduction, especially in applications where data can be compressed while retaining critical information. This analysis showed that PCA can effectively reconstruct face images, with higher numbers of components yielding higher-quality reconstructions. Using fewer components captures broad structure, while additional components add progressively finer details. This capability makes PCA suitable for tasks where storage or computational efficiency is prioritized, such as in facial recognition systems or image compression applications.

### Key Takeaways:

- PCA enables efficient data compression while retaining critical structural information.
- Optimal component selection balances data compression with the need for detail in reconstruction.
- PCA-based reconstruction quality is proportional to the number of components, with diminishing returns as more components are added.